

phylo

M. J. Wilson

September 3, 2026

Abstract

Phylogenetic inference is a discrete search whose every objective evaluation contains a continuous optimization, and the two halves do not separate: a better topology scored with badly fitted parameters looks worse than a poor one scored well. This document reports machinery for that problem, built to the standard that each claim is pinned against a source sharing no code with what it checks. The optimization interface is model-agnostic by construction — an unconstrained parameter vector, a differentiable objective, and a map back to named constrained parameters — and fits a Potts chain, a hidden Markov model, branch lengths and the general time-reversible model through the same three methods. Discrete moves sit deliberately outside it: a move changes what the parameter vector means, so it builds a new objective rather than stepping within one, and the topology search that does so needed no change to the optimizer. Felsenstein pruning agrees with brute-force marginalization to 4.0×10^{-14} relative across three backends; fitted parameters are recovered inside intervals whose coverage is measured against its nominal rate; hill climbing reaches the exhaustively enumerated optimum, and the alignment size at which it recovers the generating topology is measured against a stated bound. A reinforcement-learning estimator is checked against an exact enumeration of its own gradient, and the cheap reward a tree-searching agent would train on is validated against the maximized one it replaces — correlation 0.9568 over all 105 topologies on six taxa, at 323 times less cost. No comparison against established inference software has been run, and no trained policy proposes moves yet; Section 4.2 states what those claims would require.

Contents

1	Introduction	2
1.1	Contributions	3
1.2	What this document does not claim	3
2	Methods	3
2.1	Notation	4
2.1.1	The optimization abstraction	4
2.1.2	The phylogenetic application	4
2.2	The likelihood	4
2.2.1	Pruning	5
2.2.2	Differentiable backends	5
2.2.3	Rate variation across sites	6
2.3	Simulation	6
2.4	Tree search	6
2.4.1	The size of the problem	6
2.4.2	Move sets	6
2.4.3	Canonical serialization	7
2.4.4	Extensions to the move set	7
2.4.5	Hill climbing	8

2.4.6	Parsimony and likelihood	8
2.5	Gradients	9
2.5.1	The optimization interface	9
2.6	Reinforcement-learning formulation	10
2.6.1	Why this reward, and what it implies	10
2.6.2	Policy gradient and its baseline	10
2.6.3	Credit assignment	11
2.6.4	Status: the estimator, on a landscape that is not a tree	11
2.6.5	The phylogenetic environment, and the reward it can afford	11
2.7	Computational structure	12
3	Results	12
3.1	The simulator generates what it says it generates	13
3.2	The likelihood agrees with a source that shares no code with it	15
3.3	The optimizer recovers parameters it was not told	15
3.4	The phylogenetic fit recovers its own parameters	19
3.5	The search finds the best tree, and needs this much data to do it	19
3.6	The reward an agent can afford tracks the one it stands in for	22
4	Discussion	22
4.1	Threats to validity	22
4.2	Planned work	24
5	Conclusions	24
6	Sources	24
6.1	Infrastructure: build, structure, and speed	25
6.2	Optimization: discrete and continuous	25
6.3	Application: the science	25
A	Background material	27
A.1	NNI and SPR	27
B	The substitution model	27
B.1	Reversibility and the root distribution	27
B.2	The general time-reversible model	28
B.3	The k -state Jukes–Cantor model	28

Intended reader. Intended reader is: well-educated developer with scientific and performance-computing background, but not an expert in a given application, e.g. phylogenetics. The body stays streamlined — citing sources rather than re-deriving standard material inline — and pushes required application background into Appendix A, referenced from the point of use. It is a higher-level overview of the current best-known approach (simulation, models) in line with the roadmap, not an exhaustive record; detail, plots, and results for support are linked to dedicated supporting documents, e.g. the results of a dedicated study to find the best of a number of approaches.

1 Introduction

Inferring a phylogeny is a mixed discrete–continuous optimization problem. The discrete half is a search over tree topologies, whose count grows as $(2n - 5)!!$ and passes 10^{18} by twenty taxa. The continuous half is a fit, per candidate, of that tree’s branch lengths and substitution

parameters — and the two do not separate, because a better topology scored with badly fitted parameters looks worse than a poor one scored well. Every evaluation of the discrete objective therefore contains a continuous optimization.

Classical heuristics propose topological moves by a fixed, hand-designed rule: nearest-neighbour interchange, subtree prune-and-regraft, and their compositions, applied greedily. The question this project asks is whether a learned policy proposes better ones at a matched evaluation budget.

That question has a prerequisite the literature usually leaves implicit. A learned proposal policy is only worth building on top of machinery whose answers are known to be right, and ”known to be right” has a specific meaning here: pinned against a source that shares no code with the thing it checks. This document reports that machinery and the checks that pin it.

1.1 Contributions

1. An optimization interface that is model-agnostic by construction — an unconstrained parameter vector, a differentiable objective, and a map back to named constrained parameters — demonstrated by fitting four instances through it unchanged: a Potts chain, a hidden Markov model, branch lengths on a tree, and the general time-reversible model (Section 2.5.1).
2. The identification of discrete moves as an operation *outside* that interface: a move changes what the parameter vector means, so it constructs a new objective rather than stepping within one (Section 2.5.1). Tree search is the first consumer of that seam, and needed no change to the optimizer.
3. Felsenstein pruning in three backends, each pinned against brute-force marginalization to a relative deviation of 4.0×10^{-14} (Figure 4).
4. Parameter recovery inside confidence intervals from the observed Fisher information, with the intervals’ coverage measured against its nominal rate rather than assumed (Figures 5 and 6), and the estimability of the root branch pair established rather than asserted (Figure 7).
5. The data requirement for topology recovery, measured against this project’s own accuracy bound (Figure 11).
6. A reinforcement-learning estimator checked against an exact trajectory-enumeration oracle, and a cheap phylogenetic reward validated against the expensive one it stands in for (Figure 12).

1.2 What this document does not claim

No comparison against IQ-TREE 2, RAxML-NG, or any other established inference package has been run. Nothing here is evidence of competitiveness, and the figures that would supply it are marked as outstanding in Section 4.2. The largest problem exercised is eight taxa, bounded by the exhaustive enumeration used as the oracle; no result is reported on GPU hardware; and no trained policy proposes topological moves yet. Claims are confined to what a test pins.

2 Methods

The machinery, in the order it composes: what the symbols mean, the likelihood, the simulator that generates data with the truth retained, the search over topologies, the continuous fit inside it, and the reinforcement-learning formulation that would replace the search’s proposal rule.

2.1 Notation

Two vocabularies are kept apart, because the machinery in this document is deliberately indifferent to which application it serves. The first describes an optimization problem; the second describes a phylogeny. Nothing in the first refers to a tree.

2.1.1 The optimization abstraction

An objective is a differentiable scalar $\mathcal{L}(\boldsymbol{\theta})$ over an unconstrained vector $\boldsymbol{\theta} \in \mathbb{R}^d$, together with a map from $\boldsymbol{\theta}$ to the constrained parameters a model is actually stated in — positive rates, distributions on a simplex. Minimization is the convention throughout, so a likelihood-based objective is a negative log-likelihood.

A *discrete move* changes the structure being fitted, and therefore changes both the meaning and the dimension d of $\boldsymbol{\theta}$. It is not a step within an optimization; it constructs a new objective.

Symbol	Meaning
$\boldsymbol{\theta}, d$	unconstrained parameter vector and its dimension
$\mathcal{L}(\boldsymbol{\theta})$	scalar objective, minimized
$\nabla \mathcal{L}, \nabla^2 \mathcal{L}$	its gradient and Hessian

Table 1: Symbols for the optimization abstraction, which is shared by every application.

2.1.2 The phylogenetic application

Let $\mathcal{A}$ be the finite state space of a single character, with $k = |\mathcal{A}|$; for nucleotide data $\mathcal{A} = \{A, C, G, T\}$ and $k = 4$. The data $\mathbf{D}$ are an alignment of n taxa over m sites, so $\mathbf{D} \in \mathcal{A}^{n \times m}$, with $\mathbf{D}_{\cdot s}$ the observed column at site s .

A phylogeny is a pair $(\mathcal{T}, \boldsymbol{\tau})$: a topology $\mathcal{T}$ whose n leaves carry the taxa, and a vector $\boldsymbol{\tau}$ of non-negative branch lengths, measured in expected substitutions per site. Internal nodes represent unobserved ancestors, and ρ denotes the root.

Symbol	Meaning
$\mathcal{A}, k$	character state space and its size
n, m	number of taxa, number of sites
$\mathcal{T}$	tree topology
$\boldsymbol{\tau}$	vector of branch lengths
ρ	the root node
$\mathbf{Q}$	instantaneous rate matrix, $k \times k$
$\mathbf{P}(t)$	transition-probability matrix over a branch of length t
$\boldsymbol{\pi}$	root (and, under reversibility, stationary) distribution
$L_u(i)$	partial likelihood at node u given state i

Table 2: Symbols for the phylogenetic application. The substitution model these describe is developed in Appendix B.

2.2 The likelihood

Sites are modelled as independent given $(\mathcal{T}, \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi})$, so the log-likelihood is a sum over columns,

$$\log \mathcal{L}(\mathcal{T}, \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi} \mid \mathbf{D}) = \sum_{s=1}^m \log \Pr(\mathbf{D}_{\cdot s} \mid \mathcal{T}, \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi}). \quad (1)$$

The site term marginalizes over the unobserved states of every internal node. Written directly, that is a sum over k^{n-1} assignments — intractable beyond a handful of taxa. Felsenstein’s pruning algorithm computes it in time linear in n by exploiting the tree’s conditional independence structure [8]; it is the sum-product algorithm on a tree [14, 12].

2.2.1 Pruning

Define the partial likelihood $L_u(i)$ as the probability of the data at the leaves below node u , given that u is in state i . At a leaf with observed state a , $L_u(i) = \mathbb{1}[i = a]$. At an internal node u with children v joined by branches of length t_v ,

$$L_u(i) = \prod_{v \in \text{children}(u)} \left(\sum_{j \in \mathcal{A}} P_{ij}(t_v) L_v(j) \right), \quad (2)$$

and at the root r the site likelihood is

$$\Pr(\mathbf{D}_{\cdot s} \mid \cdot) = \sum_{i \in \mathcal{A}} \pi_i L_r(i). \quad (3)$$

Algorithm 1 Felsenstein pruning for one site

Require: topology $\mathcal{T}$, branch lengths $\boldsymbol{\tau}$, rate matrix $\mathbf{Q}$, root distribution $\boldsymbol{\pi}$, observed column $\mathbf{D}_{\cdot s}$

Ensure: site likelihood $\Pr(\mathbf{D}_{\cdot s} \mid \mathcal{T}, \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi})$

```

1: for all nodes  $u$  in post-order do
2:   if  $u$  is a leaf with observed state  $a$  then
3:      $L_u(i) \leftarrow \mathbb{1}[i = a]$  for all  $i \in \mathcal{A}$ 
4:   else
5:      $L_u(i) \leftarrow \prod_{v \in \text{children}(u)} \sum_j P_{ij}(t_v) L_v(j)$  for all  $i \in \mathcal{A}$ 
6:   end if
7: end for
8: return  $\sum_i \pi_i L_r(i)$ 
```

Algorithm 1 costs $O(nk^2)$ per site against $O(k^{n-1})$ for direct marginalization. The two agree exactly, which makes brute-force marginalization on small trees the reference implementation the pruning code is tested against.

Partial likelihoods underflow for realistic m and n , so the implementation will rescale them per node and accumulate the log of the scaling factors. Rescaling changes the arithmetic but not the result: the rescaled and direct computations must agree to within floating-point tolerance on small problems where both run.

2.2.2 Differentiable backends

Branch lengths $\boldsymbol{\tau}$ are what gradient-based fitting and tree search differentiate through, so an accelerated backend must keep them in the autodiff graph rather than baking them into topology state. The PyTorch backend (`pruning_torch.py`, float64, CPU) takes $\boldsymbol{\tau}$ as a tensor separate from the topology and is accepted when its log-likelihood agrees with the NumPy reference to $\text{atol} = 10^{-9}$ — the same bound as the brute-force check above — never relaxed to accommodate a discrepancy. Gradients obtained from `torch.autograd` are checked against central finite differences of the NumPy oracle (step $\epsilon = 10^{-6}$) to $\text{atol} = 10^{-6}$, the coarser bound set by the finite-difference truncation and rounding error at that step size, not by the autodiff computation itself.

2.2.3 Rate variation across sites

Sites evolve at different rates. The standard treatment draws a site-specific rate multiplier r from a discretized gamma distribution with mean one, and averages the site likelihood over the c rate categories,

$$\Pr(\mathbf{D}_{\cdot s} \mid \cdot) = \frac{1}{c} \sum_{g=1}^c \Pr(\mathbf{D}_{\cdot s} \mid \mathcal{T}, r_g \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi}), \quad (4)$$

multiplying the cost of a likelihood evaluation by c [8]. *Not yet implemented.*

2.3 Simulation

Simulation is how the implementation is tested: draw data from known parameters, then check that each component recovers what the generating model implies. Generation runs in pre-order, the reverse of pruning.

Algorithm 2 Simulating an alignment under $(\mathcal{T}, \boldsymbol{\tau}, \mathbf{Q}, \boldsymbol{\pi})$

Require: topology $\mathcal{T}$, branch lengths $\boldsymbol{\tau}$, rate matrix $\mathbf{Q}$, root distribution $\boldsymbol{\pi}$, sites m , seeded generator

Ensure: alignment $\mathbf{D} \in \mathcal{A}^{n \times m}$

- 1: **for** $s \leftarrow 1$ to m **do**
 - 2: draw the root state $x_r \sim \boldsymbol{\pi}$
 - 3: **for all** nodes u in pre-order, $u \neq r$ with parent p and branch t_u **do**
 - 4: draw $x_u \sim \mathbf{P}(t_u)_{x_p, \cdot}$
 - 5: **end for**
 - 6: record x_u at every leaf u as column s
 - 7: **end for**
 - 8: **return** $\mathbf{D}$
-

Every draw uses an explicitly seeded generator, so a simulated dataset is reproducible from its seed. Topology, branch lengths, k , $\boldsymbol{\pi}$, seed and site count are declared in a parameters file, and the generating truth is retained alongside the alignment, so a dataset always carries what produced it. Simulated substitution frequencies are checked against (15) within a declared Monte Carlo tolerance.

Table 4 in Section 3 lists every simulation regression fixture at the problem sizes actually run, read directly from each fixture’s yaml; Figure 2 there shows a complete generated instance at the smallest of them — the Newick string and the leading simulated sites of the alignment.

2.4 Tree search

2.4.1 The size of the problem

The number of distinct unrooted binary topologies on n taxa is the double factorial $(2n - 5)!!$ [24], which grows faster than exponentially (Table 3). Exhaustive search is impossible past a handful of taxa, and finding the optimal tree under either the parsimony or the likelihood criterion is NP-hard [8]. Practical inference is therefore heuristic: start somewhere, propose rearrangements, keep what scores better.

2.4.2 Move sets

Two rearrangement neighbourhoods define the search graph [8]: NNI and SPR, defined in Appendix A. NNI’s neighbourhood is small, cheap, and local; SPR’s is quadratic in n , so each step costs more, but its larger reach escapes local optima that trap NNI.

Taxa n	Unrooted binary topologies $(2n - 5)!!$
10	2.0×10^6
20	2.2×10^{20}
50	2.8×10^{74}

Table 3: Growth of tree space with the number of taxa.

Since neighbouring topologies differ only near the affected edges, most partial likelihoods are unchanged between a tree and its neighbour. Caching and recomputing only the invalidated path is what makes tree search affordable, and is a design requirement on the likelihood engine rather than an optimization to add later.

2.4.3 Canonical serialization

Newick notation is not a unique encoding: one topology admits many strings, differing by the order in which children are written and by where the string is rooted. A search that keys on the raw string therefore fails to recognize a tree it has already scored, and a model trained on raw strings learns to distinguish spellings that carry no phylogenetic meaning. Algorithm 3 removes both problems by fixing a representative.

Algorithm 3 Canonical form of an unrooted binary tree with labelled leaves

Require: tree $\mathcal{T}$ with leaves labelled from a totally ordered set

Ensure: a string identifying $\mathcal{T}$'s topology uniquely

```

1: root  $\mathcal{T}$  at the leaf whose label is smallest in that order
2: for all nodes  $u$  in post-order do
3:   if  $u$  is a leaf then
4:      $\kappa(u) \leftarrow \text{label}(u)$ 
5:   else
6:     sort the children  $v$  of  $u$  by  $\kappa(v)$ 
7:      $\kappa(u) \leftarrow$  the concatenation of  $\kappa(v)$  in that order
8:   end if
9: end for
10: return  $\kappa(\text{root})$ , emitting children in sorted order
```

Rooting at the smallest label is well defined for any labelled tree, and the child order is fixed bottom-up by keys that depend only on the subtree below, so the output depends on the topology alone. Hashing the keys keeps the cost at $O(n \log n)$ [6]. Branch lengths are carried separately, as a vector in canonical node order, so one topology under two length assignments shares a key — which is what makes the key usable for memoizing likelihood evaluations across a search.

Both required properties are tested rather than assumed: two trees produce the same string exactly when they are the same topology, checked by randomly re-rooting and permuting the children of one tree and against distinct trees; and parsing a canonical string reproduces the tree that produced it.

2.4.4 Extensions to the move set

Three extensions target the same bottleneck: the number of likelihood evaluations a search spends. All three are proposals. *Not yet implemented.*

Bounding whole sets of trees. Scoring every neighbour exactly wastes work when most are poor. A cheap *upper* bound on the likelihood attainable anywhere in a set of topologies discards the whole set once it falls below the best log-likelihood already achieved, in the manner of branch and bound [6]. Two sources are worth trying: relaxations of the pruning recursion (2) that replace the product over children by something cheaper and provably no smaller; and compressed summaries of the alignment, where identical site patterns already collapse and coarser encodings trade tightness for speed [2]. Soundness is the whole game — a bound that can fall below an attainable likelihood silently discards the optimum — so each candidate needs a proof, and a test that searches small trees exhaustively for a counterexample.

Multi-moves. Composing several NNI or SPR rearrangements into one action reaches topologies no single move reaches, at the cost of a much larger action space; in reinforcement-learning terms these are temporally extended actions [25]. Rather than fixing how many rearrangements a compound move contains, draw the composition during training from a Dirichlet process, so the number of distinct compound moves grows with the evidence for them [14]. The measure of success is unchanged: fewer likelihood evaluations for the same likelihood reached.

A transformer policy over Newick strings. A topology serializes to a Newick string, which a transformer can consume directly — tokenize the string, pretrain on simulated trees, then fine-tune the policy with the usual policy-gradient machinery [23, 13, 25]. The input is the canonical form of Section 2.4.3, which settles the encoding question: one topology, one string. An architecture invariant to child order and root placement buys the same property at greater cost, and without supplying the memoization key the search needs anyway. Compression applies on top, replacing recurring subtrees by dictionary symbols — an extended alphabet over the canonical string [2] — shortening the sequence the transformer attends over.

2.4.5 Hill climbing

The simplest search that uses the move sets above: score the current topology, score every neighbour, move to the best strict improvement, stop when none improves. It is the baseline the learned policy of Section 2.6 has to beat, so it is worth stating precisely rather than assuming.

Two details are not incidental. A candidate is scored by *fitting* it — the continuous parameters are re-estimated for every topology, because a topology’s likelihood is its maximized likelihood and comparing trees at shared branch lengths compares nothing. And the budget is counted in candidate fits rather than in time: fitting dominates by four orders of magnitude over generating the neighbourhood, and a wall-clock budget would make a result depend on the machine that produced it.

The discrete move is what connects this section to Section 2.5.1. A move changes the structure being fitted, so it changes what the parameter vector means and how long it is; it cannot be a step inside a fit over a fixed-length vector. It builds a new objective, and this loop is what owns that construction. The optimizer needed no knowledge of trees to serve it.

2.4.6 Parsimony and likelihood

The *large parsimony* problem asks for the topology minimizing the number of state changes needed to explain the data. We keep the search problem and replace the criterion: candidates are scored by the maximized log-likelihood of (1), with the continuous parameters $(\tau, \mathbf{Q}, \pi)$ fitted per topology. Likelihood is the more principled criterion — it models the substitution process rather than counting changes, and it is statistically consistent where parsimony can be positively misleading [8] — at the cost of a far more expensive score.

2.5 Gradients

Fitting $(\tau, \mathbf{Q}, \pi)$ for a candidate topology is a continuous optimization, and gradients make it tractable. Differentiating (11) with respect to a branch length gives

$$\frac{\partial \mathbf{P}(t)}{\partial t} = \mathbf{Q} \exp(\mathbf{Q}t) = \mathbf{Q} \mathbf{P}(t). \quad (5)$$

For a reversible $\mathbf{Q}$ with eigendecomposition $\mathbf{Q} = \mathbf{U} \Lambda \mathbf{U}^{-1}$, both the exponential and its derivative are diagonal in the eigenbasis,

$$\mathbf{P}(t) = \mathbf{U} e^{\Lambda t} \mathbf{U}^{-1}, \quad \frac{\partial \mathbf{P}(t)}{\partial t} = \mathbf{U} \Lambda e^{\Lambda t} \mathbf{U}^{-1}, \quad (6)$$

so a single decomposition serves every branch. Gradients of the log-likelihood with respect to all branch lengths follow from reverse-mode automatic differentiation through the pruning recursion (2), at a cost proportional to one forward evaluation [10, 22]. The optimizer consuming them is a standard constrained quasi-Newton or trust-region method [19]; the parameters carry constraints — non-negative branch lengths, a normalized π — so the geometry of the statistical manifold, and the natural gradient it induces, is the principled setting [1].

Analytic and automatic gradients are checked against central finite differences,

$$\frac{\partial f}{\partial \theta} \approx \frac{f(\theta + h) - f(\theta - h)}{2h}, \quad (7)$$

with h chosen to balance truncation against round-off. This check is a required test for any differentiable component, and is implemented for every objective the optimizer accepts.

2.5.1 The optimization interface

The optimizer is deliberately ignorant of what it optimizes. It takes an unconstrained parameter vector, a differentiable scalar to minimize, and a map back to the named constrained parameters; nothing in it refers to a tree, and a test asserts that the module imports nothing from the simulation, likelihood or search packages. Constraints are imposed by construction rather than by projection — a positive scale through a log map, a distribution through a softmax — so no step can leave the feasible set.

That map must also be injective. A softmax over n logits is invariant to adding a constant to all of them, so one direction of the parameter space is exactly flat; pinning the first logit removes it and makes the map a bijection onto the simplex. The same gauge appears in the Potts field of Section 3 and is fixed the same way.

Discrete moves sit outside this interface. A move changes the structure — a different topology, chain length or state count — and so changes what the parameter vector means and how long it is. It therefore constructs a *new* objective rather than stepping inside a fit over a fixed-length vector, and the loop that proposes moves owns that construction. This is what lets the tree search of Section 2.4 and the continuous fit share an optimizer without either knowing about the other.

Fitted parameters are reported with intervals from the observed Fisher information — the Hessian of the negative log-likelihood at the optimum — pushed through the constraint map by the delta method. A flat direction leaves that information singular and the interval undefined rather than merely wide, so reporting is gated on the ratio of its smallest to its largest eigenvalue; rounding leaves an exactly flat direction tiny, not zero. The gate is not hypothetical here: the two branches below a rooted binary tree’s root realize a ratio of 6.7×10^{-8} , against 6.1×10^{-2} for the same tree with that pair fitted as their sum. The intervals are asymptotic, and Section 3 reports how fast they reach their nominal coverage and shows the root invariance directly.

2.6 Reinforcement-learning formulation

Classical heuristics propose moves by a fixed, hand-designed rule. We ask whether a learned policy proposes better ones, and cast the search as a Markov decision process $(\mathcal{S}, \mathcal{A}, P, R, \gamma)$ [25]:

State. The current topology with its fitted continuous parameters, together with a summary of the alignment.

Action. A move drawn from the NNI or SPR neighbourhood of the current topology.

Transition. Applying the move and refitting the continuous parameters — deterministic given the action, up to the optimizer.

Reward. The improvement in maximized log-likelihood, $R = \log \mathcal{L}(\mathcal{T}') - \log \mathcal{L}(\mathcal{T})$.

Episode. One search trajectory from a starting tree, terminating on a step budget or when no move improves the score.

The natural performance measure is not the likelihood reached but the likelihood reached *per likelihood evaluation*: every method converges given unbounded evaluations, so the comparison against a hill-climbing baseline has to be made at equal budget.

2.6.1 Why this reward, and what it implies

The reward is a difference of log-likelihoods, which is not an arbitrary choice: it makes the undiscounted return telescope. Over a trajectory $\mathcal{T}_0, \dots, \mathcal{T}_T$,

$$\sum_{t=0}^{T-1} R_t = \sum_{t=0}^{T-1} [\log \mathcal{L}(\mathcal{T}_{t+1}) - \log \mathcal{L}(\mathcal{T}_t)] = \log \mathcal{L}(\mathcal{T}_T) - \log \mathcal{L}(\mathcal{T}_0), \quad (8)$$

so the return of an episode is exactly the total improvement it achieved, independent of the path taken to it. A policy maximizing expected return is therefore maximizing the quantity the benchmark measures, with no reward shaping to justify separately.

This forces $\gamma = 1$. Any $\gamma < 1$ breaks the telescoping and prefers improvement found early over the same improvement found late — defensible when a trajectory may not terminate, but here episodes are capped by an evaluation budget, so the return is a finite sum and needs no discounting to converge.

2.6.2 Policy gradient and its baseline

The action set is the move neighbourhood, whose size varies with n and with the move type, so the policy scores candidate moves rather than indexing a fixed action space: $\pi_{\theta}(a \mid s) \propto \exp f_{\theta}(s, a)$ over the available a . The gradient of expected return is the standard score-function estimator [25],

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t \mid s_t) (G_t - b(s_t)) \right], \quad G_t = \sum_{u=t}^{T-1} R_u, \quad (9)$$

with $b(s_t)$ a baseline that leaves the estimator unbiased — it multiplies a term of zero expectation — while reducing its variance. The variance matters more here than the bias-free property: G_t is a sum of log-likelihood differences whose scale depends on the alignment, so without a baseline the gradient magnitude varies by orders of magnitude between problems. A learned state-value estimate $\hat{v}(s_t)$ is the usual choice and is the one that makes $G_t - b(s_t)$ an advantage.

2.6.3 Credit assignment

Tree search makes the credit-assignment problem sharp. A move that unlocks a region of the space may itself lower the likelihood, with the gain arriving several moves later; by (8) the return correctly attributes that gain to the whole trajectory, but the per-step estimator in (9) must still spread it over the actions that earned it. Two consequences for an implementation. Episodes cannot be truncated arbitrarily, since a truncation that lands between the sacrifice and the payoff teaches the opposite of the truth. And a greedy baseline — hill climbing — is exactly a policy that refuses every negative R_t , so beating it requires the agent to accept some, which is the behaviour the variance-reduced estimator has to be able to learn rather than average away [25, 13].

2.6.4 Status: the estimator, on a landscape that is not a tree

The machinery above is implemented and the phylogenetic environment is not. That split is deliberate. An interface justified by one application is shaped by it, so the estimator is developed against a reference environment chosen for having nothing to do with trees: single-site flips over the same 1-D Potts chain that Section 3 fits continuously. One model, appearing once as an objective and once as a search problem, is what makes the shared abstraction a demonstration rather than a claim.

Two properties of that landscape make it a usable testbed. Its reward decomposes exactly into the two features the policy scores, so hill climbing is the weight vector proportional to $(J, 1)$ — the baseline sits *inside* the policy class, and beating it is a statement about learning rather than about two unrelated algorithms. And it is small enough to enumerate: every configuration can be scored, so “found the best one” is checkable, exactly as Section 2.4.5 enumerates topologies below eight taxa.

Enumeration also supplies the oracle the estimator itself is checked against. With a finite action set and a finite horizon the trajectory set is finite, so $J(\theta)$ is a closed form and ∇J follows by differentiating it; the sampled estimator is then an unbiasedness claim to be checked rather than cited. Measured on that oracle: the enumerated gradient agrees with central finite differences to 1.5×10^{-11} relative, and the sampled estimator with the enumerated gradient to 9.9×10^{-3} over 6000 episodes. A myopic variant — crediting each action with its own reward instead of everything that followed — is off by 71%, which is the margin that makes the check worth running.

On that landscape a learned policy beats hill climbing at a matched decision budget, reaching the enumerated optimum from 86.6% of the 81 starting configurations against greedy’s 80.2%, in 8 of 8 training seeds. The mechanism is the one Section 2.6.3 predicts: the episode ends at a local maximum, so a policy that occasionally accepts a negative R_t can route around a barrier a greedy searcher walks into. This is a harder problem than the tree fixture of Section 2.4.5, where the optimum led the runner-up by 41.6 log units and hill climbing was never challenged.

2.6.5 The phylogenetic environment, and the reward it can afford

The application instance is now built: a state is a topology, an action is a neighbour under NNI or SPR, and the reward is the improvement in log-likelihood. It lives in the search package rather than beside the estimator, for the reason the phylogenetic objective lives beside the likelihood — an agent developed against a tree is an agent shaped by trees, and the module that must stay model-agnostic may import no application code.

The reward is where the cost is. Scoring a candidate at its *maximized* branch lengths is the honest quantity and runs an optimization per candidate; an episode evaluates a whole neighbourhood at every step, so training on it is not affordable. Scoring at fixed known parameters instead is the substitution Figure 12 measures.

One detail of that substitution is forced rather than chosen. “The known parameters” do not survive a change of topology: a branch length belongs to an edge, and a different topology has different edges. The only thing that transfers is a scalar, so every candidate is scored at one fixed branch length — which makes the cheap reward a *different surface* from the fitted one rather than an approximation of it, and is why the two are compared rather than assumed interchangeable. This is the same fact Section 2.5.1 states about discrete moves, seen from the other side: a move changes what the parameter vector means.

Measuring that comparison turned up something about the fitted surface worth recording on its own. It does not totally order topologies: many candidates share a maximized log-likelihood to within the optimizer’s own convergence, because the branch that would distinguish them is fitted to zero and the tree collapses to the same polytomy — the effect Section 2.4.5’s rejected topology shows one instance of. Their relative order is therefore not a property of the model, and a rank correlation, which depends on it, is not a measurement: perturbing the fitted scores by one part in 10^9 moves Spearman’s ρ by up to 0.04. The comparison is reported as a linear correlation, which moves by the size of the perturbation instead. The cheap surface has no such degeneracy, since fixing the branch lengths at a positive value is exactly what forbids the collapse.

Not claimed here: that a learned policy beats hill climbing on trees. The 6-taxon fixture cannot support that claim in either direction — greedy already reaches the enumerated optimum from every start, as Section 2.4.5 reports, so there is nothing left to win. A fixture that separates them is a harder problem than exhaustive enumeration can referee, and that is the obstacle between here and Milestone 8 rather than any missing machinery.

Not yet implemented: PPO and a learned state-value baseline.

2.7 Computational structure

Likelihood evaluation dominates the run time: every proposed move costs at least one evaluation, and search proposes many. Its structure is favourable — sites are independent by (1), and the per-node work in (2) is a small dense matrix–vector product — so the computation is a batched, data-parallel reduction over a tree.

The project’s performance rule follows from that. Where a kernel plausibly achieves an order of magnitude over the vectorized NumPy reference at the sizes actually run, it targets the GPU through PyTorch, Triton, or JAX, which supply autodiff and batching alongside the parallelism; otherwise the work goes to the Rust backend. The threshold is a measured benchmark, not an expectation, and every accelerated kernel keeps its reference implementation as a pinned oracle. Kernel-level considerations — memory hierarchy, occupancy, parallel reduction — follow Hwu et al. [11]; the CPU-side counterparts, caching and where time actually goes, follow Bryant and O’Hallaron [4], and the backend itself is Rust [3].

Memory is the constraint to watch: partial likelihoods occupy $O(nmkc)$ floating-point values for n taxa, m sites, k states, and c rate categories, before any caching of unchanged subtrees across proposals.

3 Results

The test suite is expected to establish scientific validity, not to exercise syntax: fixtures come from component-wise simulation under known parameters, expected values come from analytic or brute-force sources, and the invariants stated above — stochasticity of $\mathbf{P}(t)$, detailed balance (12), agreement of pruning with direct marginalization, gradients against finite differences (7), re-rooting invariance — are checked directly.

Part of that suite emits the figures below. Each is regenerated on every build of this document, from the fixtures the test suite runs, and its caption is read verbatim from the file its generating

script writes. So a figure cannot drift from its description, nor either from what is tested. That holds for all of them; the paragraphs below say only why each is here, since its caption already says what it shows.

The results are grouped by the claim they support, not by the order the build emits them. Each subsection states one claim and the evidence for it.

3.1 The simulator generates what it says it generates

Everything downstream consumes simulated data, so a fault here would be invisible in every figure after it and fatal to all of them. Figure 1 opens the sequence with the object everything downstream consumes.




Figure 1: Simulated topology for 8 taxa under the Jukes-Cantor model (seed 20260904), branch lengths labelled in expected substitutions per site. The first 14 simulated sites are shown beside each leaf, in tree order, so the alignment the likelihood consumes can be read off against the topology that generated it.

Table 4 records the sizes every simulation regression test actually runs at, read from the fixtures rather than restated by hand.

Fixture	Taxa	Sites	Seed	Tolerance
<code>simulation_params.yaml</code>	4	200_000	20260902	0.01
<code>simulation_params_small_sites.yaml</code>	4	20_000	20260903	0.03
<code>simulation_params_8taxa.yaml</code>	8	200_000	20260904	0.01

Table 4: Problem-size parameters across the 3 simulation regression fixtures backing the Jukes-Cantor simulation test: taxa count, site count, seed, and the Monte Carlo tolerance each validates simulated substitution frequencies within.

Figure 2 is one complete instance at the smallest of those sizes, small enough that the alignment and the tree that produced it can be read against each other.

Figure 3 checks the simulator against eq. (15), the model it is supposed to realize. That agreement is what licenses simulated data as ground truth everywhere else here.

Figure 5 is the acceptance test for the optimizer of Section 2.5.1, on two models sharing nothing but a factorized structure: a Potts chain in an external field and a discrete hidden

(A_0.1,B_0.25,(C_0.15,D_0.4) α _0.05) ρ

	0	1	2	3	4	5	6	7	8	9
A	T	G	A	T	G	G	G	T	A	G
B	A	G	T	T	T	T	G	T	A	G
C	T	G	A	T	T	G	G	T	A	G
D	G	G	A	T	C	T	G	C	T	A

Figure 2: Worked example: 4-taxon alignment simulated under the Jukes-Cantor model (seed 20260902, 200_000 sites total), showing the topology with branch lengths – rho the root, Greek letters its internal ancestors – and the first 10 of 200_000 simulated sites.




Figure 3: Closed-form Jukes-Cantor transition probabilities for $k = 4$ (curves) against the frequencies phylo.sim.simulate produced (markers), one pair per branch of the 5-branch fixture tree. Simulated at seed 20260902 over 200000 sites. Largest deviation $1.07\text{e-}03$, within the fixture's declared Monte Carlo tolerance of $1\text{e-}02$. The simulator is validated against this analytic form, never against the likelihood code it feeds.

Markov model. Neither is phylogenetic, which is the point. An interface demonstrated on one model is shaped by that model.

3.2 The likelihood agrees with a source that shares no code with it

Three backends compute the same recursion, so agreement among them says only that they were written from the same description. What licenses the recursion is brute-force marginalization over every ancestral assignment, which shares no code with it and costs k^m for m internal nodes — the reason the check runs where it does and not at scale. Figure 4 reports it.




Figure 4: Every pruning backend against the oracle that owes it nothing. Brute-force marginalization sums over all ancestral assignments, so it shares no code with the recursion it checks, and its cost is what keeps the fixture small. Fixture: 4-state Jukes-Cantor on the 3-child-root tree, seed 20260902. The deviation is relative because the log-likelihood is a sum over sites: an absolute bound fixed at one site count does not transfer to another, which is why the x-axis spans a factor of 30. Worst deviation across every backend and size: 3.99e-14, against a float64 epsilon of 2.22e-16.

3.3 The optimizer recovers parameters it was not told

Coverage on one dataset is a draw, not a rate, so Figure 6 measures it over independent datasets across a range of sizes. Both models converge on the nominal rate from opposite sides. That is what an asymptotic approximation does; a wrong variance formula would not improve with sample size.

Figure 7 applies that optimizer to the phylogenetic case of Milestone 4. Panel (b) is why the fit has one fewer parameter than the tree has branches: moving length between the two branches below a rooted root, at fixed sum, leaves the log-likelihood unchanged to floating-point precision, while the same move between non-root siblings costs tens of log units. That is the pulley principle of Section B.1 as a statement about estimability. The pair contributes one parameter, not two, and fitting both leaves the observed information singular and every interval in panel (a) undefined.




Figure 5: Parameter recovery for the two reference instances of the optimization interface, fitted by the same model-agnostic `phylo.opt.fit`. Bars are 95 percent Wald intervals from the observed information, propagated to the reported parameters by the delta method; the dotted line is $y = x$. (a) A 3-state Potts chain, 400 chains of length 12, seed 20260902: coupling and gauge-fixed field, 4 of 4 intervals covering (1.00). (b) A 3-state, 4-symbol hidden Markov model, 600 sequences of length 15, seed 20260903: initial, transition and emission probabilities after aligning the hidden-state permutation the likelihood is invariant to, 23 of 24 intervals covering (0.96). Coverage on a single dataset is a draw, not a rate: the nominal 0.95 is checked over independent replicates by the fitting regression suite, and against sample size in the companion coverage figure.




Figure 6: Realized coverage of the 95 percent Wald intervals built from the observed information, against the amount of data each fit sees. Bars are binomial standard errors on the realized fraction. Both instances converge on the nominal rate and neither sits on it at the smallest sample, which is what an asymptotic approximation does and a wrong formula does not. They approach from opposite sides. The Potts chain over-covers and comes down: $158/160 = 0.988$ at 100 chains, $96/100 = 0.960$ at 1600. The hidden Markov model under-covers and comes up: $45/48 = 0.938$ at 150 sequences, $91/96 = 0.948$ at 2400. The under-coverage has two identifiable sources absent from the Potts chain – some emission probabilities are fitted near zero, where a Wald interval on the log scale is a poor approximation, and aligning the hidden-state permutation to the truth is a post-selection step that costs a little coverage. Seeds are fixed, so every point is reproducible rather than redrawn. At the smallest hidden Markov size 6 of 18 fits reached the boundary of the parameter space – an emission probability estimated at zero – where the observed information is singular and there is no interval to report. Those fits are excluded from the counts above and stated here rather than dropped silently, since excluding them unannounced would select for the well-behaved samples.




Figure 7: Branch-length fitting by autodiff, using the same model-agnostic optimizer as the Potts and hidden Markov instances. (a) Every estimable branch length against the value that generated the alignment, with 95 percent Wald intervals from the observed information; markers are open where an interval misses. Fixtures: 3-child root, seed 20260903, and 2-child root, seed 20260904, both scored at 5000 sites. (b) Why some branch lengths are not estimable at all. Moving mass between the two branches below a rooted binary tree’s root, at fixed sum, changes the log-likelihood not at all: the log-likelihood is bit-identical across a 9 to 1 range. The same move between two non-root siblings changes it by up to 90.1. This is the pulley principle: under a reversible model the likelihood does not depend on where the root sits along the branch it subdivides. Only the pair’s sum is estimable, so the objective fits it as one parameter; fitting both would leave the observed information singular and every interval in panel (a) undefined.

3.4 The phylogenetic fit recovers its own parameters

Figure 8 fits the substitution model itself — the nine free parameters of Section B.2 alongside the branch lengths, through the same interface. Panel (b) is the half that is easy to omit, and the caption says why.




Figure 8: Fitting the general time-reversible model: five free exchangeabilities (the sixth is pinned at 1) and four stationary frequencies, alongside the branch lengths, by the same model-agnostic optimizer used for the Potts and hidden Markov instances. Bars are 95 percent Wald intervals from the observed information; the dotted line is $y = x$; markers are open where an interval misses. Topology: the 2-child-root fixture, seed 20260904, 20000 sites. (a) Data simulated under an asymmetric truth, 8 of 9 intervals covering. (b) Data simulated under Jukes-Cantor, where every exchangeability is 1 and every frequency $1/4$: 9 of 9 covering. The second panel is the half that is easy to forget. A model flexible enough to fit panel (a) is only useful if it also declines to invent structure that is not there, and Jukes-Cantor is exactly the point of the parameter space where the general model could most easily do so. Coverage on a single dataset is a draw and not a rate: the nominal 0.95 is checked over independent replicates by the release-gated regression suite.

Figure 9 needs both its panels. A trajectory alone shows a search improving without saying whether it stopped anywhere good; the landscape alone shows the problem without saying whether the search solved it. Below eight taxa the enumeration in panel (b) is affordable, and it is what turns “the search found the best tree” from an assertion into a checkable claim — the only independent oracle this project has for the discrete half of the problem.

3.5 The search finds the best tree, and needs this much data to do it

Figure 10 shows what that discrimination amounts to. The search recovers the generating topology here, so inferred against true would draw the same tree twice; the informative comparison is the




Figure 9: Hill climbing over tree topologies, scoring every candidate with the same model-agnostic optimizer used for the Potts and hidden Markov instances. Fixture: 4-state Jukes-Cantor, 1500 sites, seed 20260905. (a) Log-likelihood against candidate fits spent, from a randomly drawn starting topology; the dotted line is the generating tree’s own score. (b) All 105 unrooted topologies on this leaf set, scored and ranked, with the search’s endpoint dashed. Below 8 taxa this enumeration is affordable, which is what makes the phrase found the best tree a checkable claim rather than an assertion. The optimum leads the runner-up by 41.6 log units here, so this problem does not yet separate the two move sets: both reach it, NNI in fewer fits.

runner-up it ruled out. They differ by one split, and the rejected tree fits the internal branch that would create the wrong grouping at essentially zero length — which is what rejecting a topology looks like from inside the continuous fit.




Figure 10: The topology hill climbing selected, beside the highest-scoring one it rejected, each with its own branch lengths fitted. Fixture: 4-state Jukes-Cantor over 1500 sites, seed 20260905. (a) scores -8775.5 and is the generating topology; (b) scores -8817.1, a difference of 41.6 log units. They differ by one split: C D are a group in one and not in the other. On this fixture the search recovers the generating tree, so a drawing of found against true would show the same tree twice; the runner-up is what the likelihood actually had to discriminate against. Note its shortest internal branch, fitted at 0.000: a wrong grouping is generally only supportable by collapsing the edge that would create it.

Figure 12 is the one an agent depends on. An episode evaluates a whole neighbourhood at every step, so a reward costing one optimization per candidate puts training out of reach; scoring at fixed known parameters instead is cheap enough, and the question is whether it is the same problem. Panel (b) exists because the cheap surface has exactly one free parameter, and a result that held only at the value we happened to pick would not be a result.

Both of those run on one alignment. Whether the search recovers the generating tree at all depends on whether the alignment carries the signal to distinguish it, and `ROADMAP.md` sets a bound on the resulting error — a normalized Robinson–Foulds distance of 0.05 — that nothing had measured. Figure 11 measures it against the site count, which is the quantity that decides the answer.




Figure 11: How much alignment the search needs to recover the tree that generated it. Each marker is one independent alignment drawn from the 4-state Jukes-Cantor fixture (seed 20260905, 8 replicates per size), inferred by NNI hill climbing and scored against the generating topology by normalized Robinson-Foulds distance – 0 for the same tree, 1 for no internal split in common. The dotted line is the accuracy requirement this project set itself, 0.05. It is met from 125 sites upward: 8 of 8 replicates recover the topology exactly at 2000 sites, against 5 of 8 at 60. The distance is normalized by internal splits, so the bound means the same thing at any taxon count.

3.6 The reward an agent can afford tracks the one it stands in for

4 Discussion

Two results are worth separating from the rest, because they were not what was expected.

The first is that the abstraction held. An optimization interface written without reference to any model was expected to need widening once a real application arrived; it did not. Branch lengths, exchangeabilities and stationary frequencies fit through the same three methods that fit a Potts chain, and the topology search that consumes them required no change to the optimizer at all. The seam that made that work is the one in Section 2.5.1: a discrete move changes what the parameter vector means, so it builds a new objective instead of stepping inside one. Stating that as an exclusion, rather than absorbing moves into the optimizer, is what kept the interface model-agnostic.

The second is that the fitted likelihood does not totally order topologies. Fitting each of the 105 unrooted topologies on six taxa produces only 52 distinct maxima: where the branch that would distinguish two trees is fitted to zero, both collapse to the same polytomy and share a likelihood exactly. This is a property of the objective, not of the optimizer, and it has a practical consequence — any statistic over topology rankings is discontinuous in the fitted values, so a rank correlation computed on them is not reproducible across machines. Section 3.6 reports a linear correlation for that reason.

4.1 Threats to validity

The exhaustive enumeration that makes “found the best tree” checkable also bounds where it can be checked: $(2n - 5)!!$ passes ten thousand at eight taxa. Every search result here is therefore established at a size where hill climbing is not challenged — on the six-taxon fixture




Figure 12: The reward an agent can afford, against the one it stands in for. Every one of the 105 unrooted topologies on this leaf set, scored twice: at fixed known parameters, and at parameters maximized per topology by the same optimizer the Potts and hidden Markov instances use. Fixture: 4-state Jukes-Cantor, 1500 sites, seed 20260905; the fixed branch length is the generating tree’s mean, 0.1644. (a) The two surfaces, correlation 0.9568; both score the generating topology highest, so the cheap reward and the expensive one agree on the answer while disagreeing in detail. The correlation is linear rather than rank-based because the fitted surface does not totally order topologies: several of its optima agree to within the optimizer’s own convergence, so their relative order is not a property of the model. (b) The cheap surface has one free parameter, and the conclusion does not rest on it: across a 50-fold range of fixed branch length the two surfaces pick the same best topology in 7 of 7 cases, with the correlation never falling below 0.8719. The substitution is what makes training affordable at all: the fitted reward runs a full optimization per candidate, the known one a single pruning pass.

the optimum leads the runner-up by 41.6 log units and greedy reaches it from every start. That the searches succeed is evidence the machinery is correct, and no evidence at all about the hard instances the project exists to attack.

The simulated data is generated by the same models the likelihood assumes, so model misspecification is untested by construction. Real alignments violate those assumptions, and Milestone 9 is where that is confronted.

4.2 Planned work

The claims this document cannot yet make, and the figures that would make them:

Competitiveness. Maximized log-likelihood against evaluation budget, for the learned policy, hill climbing, IQ-TREE 2 and RAxML-NG on shared seeds. Requires those packages in the environment and the benchmark harness of Milestone 7. Until it exists, no statement in this document compares the method to established software.

A trained phylogenetic agent. The learning curve of a policy over tree moves. The obstacle is a fixture rather than machinery: separating a policy from greedy needs a problem harder than exhaustive enumeration can referee, so the oracle that validates the search cannot validate the agent that replaces it.

Scale. Runtime and memory against n and L across the range `ROADMAP.md` specifies. Memory is deterministic and can be reported here; runtime cannot, since a figure in this document is regenerated by continuous integration on hardware that varies, and a timing committed into it would neither reproduce nor mean anything. Those numbers live in the benchmark suite.

Rate variation across sites, and GPU dispatch. Both are specified in Section 2.2.3 and Section 2.2.2 and neither is built.

5 Conclusions

A model-agnostic optimization interface serves four instances unchanged, and excluding discrete moves from it is what makes that possible rather than a limitation of it. Felsenstein pruning agrees with brute-force marginalization to 4.0×10^{-14} relative across three backends. Parameters are recovered inside intervals whose coverage is measured rather than assumed. Hill climbing reaches the exhaustively enumerated optimum, and the alignment size at which it does so is measured against this project’s own accuracy bound. A reinforcement-learning estimator is checked against an exact enumeration of its own gradient, and the cheap reward a phylogenetic agent would train on is validated against the expensive one it replaces.

What remains is the part the project is named for: a policy that proposes better moves than the hand-designed rule, demonstrated against established software on problems large enough for the comparison to matter. The machinery for it is built and pinned. The evidence is not.

6 Sources

The texts this project works from, grouped as `CLAUDE.md` groups them — a routing table keyed on what each informs, rather than a reading list. Where an algorithm we implement appears in one of them, this document follows its formulation and notation and says where it deviates.

6.1 Infrastructure: build, structure, and speed

Software craft. Martin [15] on the structure that keeps a change reviewable; Blandy et al. [3] on ownership, lifetimes, and the FFI boundary the Rust backend crosses.

Systems and hardware. Bryant and O’Hallaron [4] on the memory hierarchy and where CPU time goes; Hwu et al. [11] on GPU kernels, occupancy, and parallel reduction and scan.

6.2 Optimization: discrete and continuous

Algorithms and discrete mathematics. Cormen et al. [6] on data structures, traversal, complexity, and branch and bound; Rosen [24] on the counting and graph theory behind tree space.

Numerical optimization. Nocedal and Wright [19] on line search, trust region, and quasi-Newton methods for the continuous parameters.

Probabilistic inference. MacKay [14] on inference, message passing, and Monte Carlo; Koller and Friedman [12] on factor graphs and the exact-inference machinery that (2) instantiates; Frey [9] on belief propagation read across inference and coding; Ortega [20] on spectral methods for signals indexed by graph vertices.

Statistical physics. Mézard and Montanari [16] on inference over sparse graphs and phase transitions in hard combinatorial problems — tree search being one; Newman and Barkema [17] on sampling, equilibration, and error estimates from correlated samples, which the benchmark harness needs to rank variants rather than noise.

Information geometry. Amari [1] on the statistical manifold, the Fisher metric, and natural gradient.

Learning and reinforcement learning. Goodfellow et al. [10] and Prince [22] on architectures, optimization, and automatic differentiation; Sutton and Barto [25] on MDPs, value and policy methods, and temporally extended actions; Lapan [13] on implementation practice; Raschka [23] on transformer internals and tokenization, for the policy over canonical Newick strings proposed in Section 2.4.4.

6.3 Application: the science

Phylogenetics. Felsenstein [8] on substitution models, pruning, tree search, and the parsimony and likelihood criteria; Durbin et al. [7] on probabilistic models of sequences; Compeau and Pevzner [5] on the algorithmic view of alignment and phylogeny construction; Pachter and Sturmfels [21] on the algebraic structure of phylogenetic models and their invariants, which offer topology tests that avoid likelihood optimization altogether.

Information and quantum. Blahut [2] on algebraic coding, which is the background for the compressed tree encodings in Section 2.4.4; Nielsen and Chuang [18] on quantum information and algorithms, background only — nothing here targets quantum hardware.

References

- [1] Shun-ichi Amari. *Information Geometry and Its Applications*. Springer, Tokyo, 2016.
- [2] Richard E. Blahut. *Algebraic Codes for Data Transmission*. Cambridge University Press, Cambridge, 2003.

- [3] Jim Blandy, Jason Orendorff, and Leonora F. S. Tindall. *Programming Rust: Fast, Safe Systems Development*. O'Reilly Media, 2nd edition, 2021.
- [4] Randal E. Bryant and David R. O'Hallaron. *Computer Systems: A Programmer's Perspective*. Pearson, 3rd edition, 2015.
- [5] Phillip Compeau and Pavel Pevzner. *Bioinformatics Algorithms: An Active Learning Approach*. Active Learning Publishers, 2015.
- [6] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms*. MIT Press, Cambridge, Massachusetts, 4th edition, 2022.
- [7] Richard Durbin, Sean R. Eddy, Anders Krogh, and Graeme Mitchison. *Biological Sequence Analysis: Probabilistic Models of Proteins and Nucleic Acids*. Cambridge University Press, Cambridge, 1998.
- [8] Joseph Felsenstein. *Inferring Phylogenies*. Sinauer Associates, Sunderland, Massachusetts, 2004.
- [9] Brendan J. Frey. *Graphical Models for Machine Learning and Digital Communication*. MIT Press, Cambridge, Massachusetts, 1998.
- [10] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, Cambridge, Massachusetts, 2016.
- [11] Wen-mei W. Hwu, David B. Kirk, and Izzat El Hajj. *Programming Massively Parallel Processors: A Hands-on Approach*. Morgan Kaufmann, 4th edition, 2022.
- [12] Daphne Koller and Nir Friedman. *Probabilistic Graphical Models: Principles and Techniques*. MIT Press, Cambridge, Massachusetts, 2009.
- [13] Maxim Lapan. *Deep Reinforcement Learning Hands-On*. Packt Publishing, 2nd edition, 2020.
- [14] David J. C. MacKay. *Information Theory, Inference, and Learning Algorithms*. Cambridge University Press, Cambridge, 2003.
- [15] Robert C. Martin. *Clean Code: A Handbook of Agile Software Craftsmanship*. Prentice Hall, 2008.
- [16] Marc Mézard and Andrea Montanari. *Information, Physics, and Computation*. Oxford University Press, Oxford, 2009.
- [17] M. E. J. Newman and G. T. Barkema. *Monte Carlo Methods in Statistical Physics*. Oxford University Press, Oxford, 1999.
- [18] Michael A. Nielsen and Isaac L. Chuang. *Quantum Computation and Quantum Information*. Cambridge University Press, Cambridge, 10th anniversary edition, 2010.
- [19] Jorge Nocedal and Stephen J. Wright. *Numerical Optimization*. Springer, 2nd edition, 2006.
- [20] Antonio Ortega. *Introduction to Graph Signal Processing*. Cambridge University Press, Cambridge, 2022.
- [21] Lior Pachter and Bernd Sturmfels, editors. *Algebraic Statistics for Computational Biology*. Cambridge University Press, Cambridge, 2005.
- [22] Simon J. D. Prince. *Understanding Deep Learning*. MIT Press, Cambridge, Massachusetts, 2023.

- [23] Sebastian Raschka. *Build a Large Language Model (From Scratch)*. Manning Publications, 2024.
- [24] Kenneth H. Rosen. *Discrete Mathematics and Its Applications*. McGraw-Hill, 8th edition, 2019.
- [25] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, Cambridge, Massachusetts, 2nd edition, 2018.

A Background material

Standard algorithms that are not specific to `phylo`, referenced from the main text rather than re-derived there (see the intended-reader note after the table of contents).

A.1 NNI and SPR

Two rearrangement neighbourhoods define the tree-search graph used in Section 2.4.4 and throughout [8]:

NNI (nearest-neighbour interchange). Each internal edge of an unrooted binary tree separates four subtrees, which can be rejoined in three ways; two are new topologies. With $n - 3$ internal edges, the NNI neighbourhood contains $2(n - 3)$ topologies — small, cheap, and local.

SPR (subtree prune and regraft). Detach a subtree and reattach it at another edge. The neighbourhood is quadratic in n , so each step costs more, but the larger reach escapes local optima that trap NNI.

B The substitution model

Character evolution along a branch is a continuous-time Markov chain on $\mathcal{A}$ with rate matrix $\mathbf{Q}$, whose off-diagonal entries $q_{ij} \geq 0$ give the instantaneous rate of substitution from state i to state j and whose rows sum to zero,

$$q_{ii} = - \sum_{j \neq i} q_{ij}. \quad (10)$$

Transition probabilities over a branch of length t follow from the matrix exponential,

$$\mathbf{P}(t) = \exp(\mathbf{Q}t), \quad P_{ij}(t) = \Pr(X_t = j \mid X_0 = i). \quad (11)$$

$\mathbf{P}(t)$ is a stochastic matrix: $P_{ij}(t) \geq 0$ and each row sums to one. Both properties are invariants the implementation is tested against.

B.1 Reversibility and the root distribution

The models used here are time-reversible with respect to a distribution $\boldsymbol{\pi}$, meaning they satisfy detailed balance,

$$\pi_i q_{ij} = \pi_j q_{ji} \quad \text{for all } i, j. \quad (12)$$

Reversibility has two consequences we rely on. First, $\boldsymbol{\pi}$ is the stationary distribution of the chain, which makes it the natural choice for the distribution of states at the root. Second, the likelihood is invariant to where the root is placed on the tree — Felsenstein’s *pulley principle*

[8] — so an unrooted topology suffices for inference, and re-rooting is a test the implementation must pass rather than a modelling decision.

Every reversible rate matrix factorizes in the same way, which is the general time-reversible (GTR) model of Section B.2, eq. (14); the familiar JC69, K80 and HKY85 models are constrained special cases [8]. Branch lengths measure expected substitutions per site, which fixes the scale of $\mathbf{Q}$ through the normalization

$$-\sum_i \pi_i q_{ii} = 1. \quad (13)$$

Without (13), $\mathbf{Q}$ and $\boldsymbol{\tau}$ trade off against each other and neither is identifiable.

B.2 The general time-reversible model

Jukes–Cantor has no free rate parameters, so there is nothing in $\mathbf{Q}$ or $\boldsymbol{\pi}$ to estimate under it. The general reversible model supplies some:

$$\mathbf{Q}_{ij} = s_{ij} \pi_j \quad (i \neq j), \quad \mathbf{Q}_{ii} = -\sum_{j \neq i} \mathbf{Q}_{ij}, \quad (14)$$

with s symmetric — the *exchangeabilities* — and $\boldsymbol{\pi}$ the stationary distribution. Detailed balance, $\pi_i \mathbf{Q}_{ij} = \pi_j \mathbf{Q}_{ji}$, follows immediately from the symmetry of s , so (14) is reversible by construction rather than by assumption, and the eigendecomposition (6) applies exactly.

Two normalizations are gauges rather than conventions, and a third appears once $\boldsymbol{\pi}$ is fitted. Each removes a direction along which the likelihood is exactly flat; a flat direction makes the observed information singular, so without them no parameter of the model has a confidence interval at all. First, $\mathbf{Q}$ and $c\mathbf{Q}$ describe the same process at branch lengths t and t/c , so $\mathbf{Q}$ is scaled to one expected substitution per unit time, $-\sum_i \pi_i \mathbf{Q}_{ii} = 1$; this is what makes a branch length mean “expected substitutions per site”. Second, the same freedom reappears in s , since scaling every entry is undone by that normalization, so one entry is held at 1. Third, $\boldsymbol{\pi}$ is optimized through a softmax whose first logit is pinned, since adding a constant to all of them changes nothing.

Setting every s_{ij} equal and $\boldsymbol{\pi}$ uniform must reproduce the Jukes–Cantor rate matrix exactly, which is how the implementation is validated — against a model this document has already pinned against its own closed form, rather than against a restatement of (14).

B.3 The k -state Jukes–Cantor model

The most constrained case is worth stating in closed form, because it gives the simulator an oracle that owes nothing to our likelihood code. Take all exchange rates equal and $\boldsymbol{\pi}$ uniform on k states. Under the normalization (13), the transition probabilities over a branch of length t are

$$P_{ii}(t) = \frac{1}{k} + \frac{k-1}{k} e^{-kt/(k-1)}, \quad P_{ij}(t) = \frac{1}{k} - \frac{1}{k} e^{-kt/(k-1)} \quad (i \neq j). \quad (15)$$

At $t = 0$ this gives the identity, and as $t \rightarrow \infty$ every entry tends to $1/k$, the stationary distribution. For $k = 4$ it is JC69 [8]. Simulated substitution frequencies are checked against (15) within Monte Carlo error, which tests the simulator independently of the pruning implementation — if both were validated against each other, a shared error would pass unnoticed.